

AgiSphere-CORE: A Multi-Dimensional Co-Optimization Framework for Autonomous Prescriptive Maintenance in Smart Manufacturing

Pinghao Luo

University of Southern California

Zhiguo Tao*

University of Nottingham

Abstract

Intelligent workflows for smart manufacturing reliability and maintenance (RxM) are transitioning from static, sequential multi-agent systems toward autonomous, dynamic decision-making. However, a critical gap persists between agent coordination and the holistic co-optimization required across predictive performance, computational resources, and domain-specific business logic. To address this, we propose **AgiSphere-CORE**, a novel framework centered on a **Co-Optimization Core (COC)**. The COC maintains a unified *Collaborative State Vector* and solves a constrained optimization problem to maximize the *Collaborative Utility* J_{co} , dynamically orchestrating specialized agents for resources, models, and logic. This enables a paradigm shift from sequential to co-optimization-driven autonomy. Extensive evaluation on multiple manufacturing tasks—including classification, regression, and anomaly detection—demonstrates that AgiSphere-CORE consistently outperforms baselines, achieving superior predictive performance (e.g., 97.3% accuracy, 0.902 F1-score), improved resource efficiency (15% faster workflow time), and higher estimated domain cost savings (\$148.7K). Ablation studies confirm the critical role of the COC and the adaptive learning mechanism. A detailed case study further illustrates the framework’s transparent, utility-driven trade-offs. We conclude that AgiSphere-CORE provides a robust, holistic foundation for next-generation autonomous RxM systems in smart manufacturing.

1 Introduction

Manufacturing is undergoing a transformative shift toward data-driven, resilient, and autonomous operations, driven by the convergence of Cyber-Physical Systems (CPS), the Industrial Internet of Things (IIoT), and advanced analytics. Inspired

by the foundational work on agricultural automation and sustainable systems (Jiang et al., 2025a,b), this evolution has elevated predictive and prescriptive maintenance (RxM) from reactive diagnostics to a cornerstone of operational resilience. However, modern manufacturing environments pose significant challenges due to their dynamic nature, data intensity, and inherent uncertainty. Traditional centralized or rule-based systems often struggle to meet demands for scalability, adaptability, and real-time decision-making (Wu et al., 2022, 2024). While predictive maintenance (PdM) effectively forecasts failures, contemporary manufacturing requires systems that not only anticipate issues but also recommend or execute optimal corrective actions. Building upon the adaptive federated learning paradigm established by (Wu et al., 2022) and extending the multi-objective optimization framework of (Wu et al., 2024), this necessity calls for distributed intelligence that seamlessly integrates reasoning with coordinated execution.

Multi-Agent Systems (MAS) offer a promising paradigm for distributed intelligence in industrial settings by enabling autonomous entities to collaborate toward common goals (Li, 2026; Yang et al., 2025; Firdaus et al., 2024). Yet, their effectiveness is frequently limited by agents with shallow reasoning depth and a reliance on predefined rules, which reduces adaptability to unforeseen events (Guan et al., 2024; Ullah et al., 2021; Wang et al., 2017). Conversely, recent advances in Agentic AI, powered by foundation models, provide sophisticated reasoning, planning, and tool-use capabilities (Ravikumar et al., 2022; Ismail et al., 2021; Adade et al., 2025). However, systems built solely on large language models (LLMs) face computational, latency, and privacy challenges, while small language models (SLMs), though efficient and privacy-preserving, exhibit more constrained reasoning capacity (Wu et al., 2023; Cheng et al., 2023). A critical gap thus emerges: tradi-

*Corresponding author: ssyzt3@outlook.com.

tional MAS lack adaptive reasoning, while agentic AI systems often rely on single-agent architectures, limiting scalability and resilience. Our work significantly outperforms these baseline approaches (Jiang et al., 2025b,a) by achieving superior integration of distributed intelligence (Shabbir et al., 2020; Rasool et al., 2020).

To address this gap, inspired by recent advances in agentic AI architectures (Wu et al., 2022, 2024; Li, 2026) and building upon multi-agent coordination frameworks (Cayambe et al., 2023; Liu et al., 2021; Song et al., 2026b), we propose **AgiSphere-CORE** (Co-Optimization Framework for Autonomous RxM), a novel hybrid architecture that unifies MAS coordination principles with the reasoning-rich autonomy of agentic AI. Our core contribution is a central **Co-Optimization Core (COC)** that performs continuous multi-dimensional optimization across performance, resource, and domain-logic states, orchestrating three specialized agents (Resource, Model, Logic) through a unified collaborative utility function. This framework integrates SLMs for lightweight, local reasoning and strategically employs LLMs for high-level strategy initialization, thereby enabling adaptive, goal-directed collaboration (Zhu et al., 2024; Kong et al., 2024; Lin et al., 2019). Experimental validation across diverse manufacturing tasks demonstrates that *AgiSphere-CORE* consistently achieves superior or competitive predictive performance while delivering significant gains in resource efficiency and domain-cost savings, outperforming the baselines established by (Jiang et al., 2025a; Wu et al., 2022) (Issaka et al., 2018; Xiang et al., 2021). The remainder of this paper is structured as follows: Section 2 details our proposed methodology, Section 3 presents experimental results and analysis, and Section 4 concludes the work.

2 Related Work

Research on distributed decision-making using Multi-Agent Systems (MAS) in manufacturing has demonstrated benefits but often suffers from rigid reasoning mechanisms (Zhang et al., 2023; Jiang et al., 2024; Lu et al., 2021). While Agentic AI, initially explored in general-purpose, single-agent prototypes based on Large Language Models (LLMs), is rapidly expanding into industrial contexts such as manufacturing and logistics (El-Mesery et al., 2024; Fayek et al., 2023), robust applications in-




Figure 1: Paradigm shift from sequential to co-optimization-driven autonomy. Traditional systems use isolated agents with sequential hand-offs, whereas AgiSphere-CORE uses a central co-optimization core to coordinate specialized agents through bidirectional information exchange.

creasingly rely on MAS and tool-calling architectures (Su et al., 2025) to achieve the necessary scalability, specialization, and reliability for enterprise workflows (Qenawy et al., 2024; Zeng et al., 2022). In smart manufacturing, advanced Machine Learning (ML) and Deep Learning (DL) methods for reliability and maintenance (RxM) show promise but often lack generalization, interpretability, and scalability in dynamic production environments (Jia et al., 2018; Qiu et al., 2017). Consequently, hybrid approaches are emerging, but the full integration of agentic AI with MAS for RxM remains nascent (Solairaj et al., 2021; Sun et al., 2019).

2.1 From Predictive to Prescriptive Maintenance

Predictive Maintenance (PdM) is a central goal for Smart Manufacturing Systems (SMS), though its adoption faces barriers related to cost, data maturity, and organizational readiness (Feng et al., 2024; Yuan et al., 2016). The field is evolving towards autonomous frameworks that integrate prediction, diagnosis, and decision-making within a closed loop, aligning with resilient and self-adaptive production objectives (Zhao et al., 2018; Cui et al., 2025). Driven by sensor data and the Industrial Internet of Things (IIoT), PdM applications span health monitoring, fault detection, and production optimization (Sun et al., 2023; Tang et al., 2024).

Methodologically, classical time-series models (e.g., ARIMA) and ML algorithms (e.g., SVM, Random Forests, XGBoost) have been widely ap-

plied for forecasting and fault classification (Wang et al., 2025a; Yu et al., 2025b). More recently, deep learning architectures—including Convolutional Neural Networks (CNNs) for feature extraction, Recurrent Neural Networks (RNNs) for sequential modeling, and autoencoders for anomaly detection—have become prevalent, often evaluated on benchmark datasets (Yu et al., 2025a,c).

Despite progress, current PdM faces critical challenges: scalability bottlenecks from centralized high-volume data streams; limited adaptability under distributional shifts; poor interpretability of complex models; and latency or privacy concerns (Jiang et al., 2025b; Wu et al., 2022). Active research addresses these via distributed or federated learning, domain adaptation, and explainable AI (XAI) (Wu et al., 2024; Jiang et al., 2025a; Hsieh et al., 2024).

Building on this, the emerging paradigm of Prescriptive Maintenance (RxM) aims to close the loop between monitoring, prediction, and action by determining optimal maintenance interventions (Li, 2026; Yang et al., 2025). It integrates diagnostic reasoning, optimization, and decision support, leveraging ML, reinforcement learning, and systems integration (Firdaus et al., 2024; Guan et al., 2024). However, large-scale implementation is constrained by data heterogeneity, validation challenges in safety-critical settings, and the need for transparent decision logic (Ullah et al., 2021; Wang et al., 2017).

2.2 Multi-Agent Systems in Manufacturing

Research on MAS in manufacturing, motivated by the need for distributed, adaptive decision-making, dates back decades (Ravikumar et al., 2022; Ismail et al., 2021). Early work focused on production scheduling and resource allocation, demonstrating improved scalability and robustness to disturbances (Adade et al., 2025; Wu et al., 2023). Applications later expanded into maintenance and diagnostics, using hierarchical or cooperative agent layers and protocols for distributed negotiation (Cheng et al., 2023; Shabbir et al., 2020).

Recent extensions include decentralized MAS for joint scheduling and maintenance planning, ontology-based frameworks for improved knowledge representation and coordination, and the integration of Multi-Agent Reinforcement Learning (MARL) for adaptive control under uncertainty (Rasool et al., 2020; Cayambe et al., 2023; Liu et al., 2021). More recently, researchers have explored

augmenting MAS with LLMs to overcome rule-based rigidity, enhancing human-machine collaboration and system resilience (Zhu et al., 2024; Kong et al., 2024). This convergence is noted as a key enabler of next-generation industrial intelligence (Lin et al., 2019; Issaka et al., 2018). In maintenance, AI-enabled MAS frameworks have been proposed for closed-loop architectures integrating fault detection, diagnosis, and simulation (Xiang et al., 2021; Zhang et al., 2023).

Collectively, MAS provides a robust foundation for distributed decision-making (Jiang et al., 2024; Lu et al., 2021). However, their reliance on predefined reasoning limits adaptability, creating an opportunity for integration with the dynamic reasoning capabilities of agentic AI (El-Mesery et al., 2024; Fayek et al., 2023).

2.3 Agentic AI and Autonomous Agents

Agentic AI represents a shift from traditional MAS agents by combining LLMs with planning, tool use (Song et al., 2026a), memory, and self-reflection mechanisms, enabling flexible interpretation of multimodal inputs and context-aware autonomous action (Qenawy et al., 2024; Zeng et al., 2022; Jia et al., 2018). Frameworks such as CrewAI, LangGraph, and AutoGen have demonstrated capabilities in software development and web-based task automation using plan-act-reflect loops (Qiu et al., 2017; Solairaj et al., 2021). Their natural language interface supports intuitive human-agent collaboration (Sun et al., 2019; Feng et al., 2024).

However, most applications have been non-industrial, focusing on information-centric tasks like retrieval or report generation (Yuan et al., 2016; Zhao et al., 2018). These domains differ significantly from manufacturing, which imposes stricter requirements for safety, reliability, latency, and robustness (Cui et al., 2025; Sun et al., 2023). Furthermore, current work emphasizes single-agent autonomy, raising scalability questions for environments like predictive maintenance that require simultaneous management of multiple assets and data streams (Tang et al., 2024; Wang et al., 2025a). Extending agentic AI principles to multi-agent industrial ecosystems thus remains a promising, unexplored direction for enabling autonomous maintenance (Yu et al., 2025b,a).

2.4 Hybrid and Emerging Approaches

Recent hybrid approaches integrate MAS with ML techniques (Yu et al., 2025c; Jiang et al., 2025b). For example, MAS frameworks have been augmented with reinforcement learning agents for adaptive decision rules or with embedded anomaly detection models, aiming to combine MAS’s distributed scalability with ML’s adaptability (Jiang et al., 2025a; Wu et al., 2022). Similarly, digital twin-based predictive maintenance has gained momentum, providing real-time simulations for diagnosis (Wu et al., 2024; Li, 2026). However, most digital twin solutions remain centralized and lack the conversational, reasoning-driven intelligence characteristic of agentic AI (Yang et al., 2025; Firdaus et al., 2024).

To date, the integration of LLM-powered agentic AI into MAS architectures remains unexplored in industrial contexts (Guan et al., 2024; Ullah et al., 2021). Such integration holds the potential to overcome the rigidity of rule-based MAS while avoiding the scalability bottlenecks of single-agent agentic AI, allowing each agent to leverage LLM-based reasoning for autonomy while MAS-style coordination ensures system-wide robustness (Wang et al., 2017; Ravikumar et al., 2022; Ismail et al., 2021).

3 Method

We propose **AgiSphere-CORE** (Co-Optimization Framework for Autonomous RxM), a novel multi-agent system that re-architects the intelligent workflow for smart manufacturing RxM. Departing from prior hybrid architectures that sequence specialized agents, our framework centers on a **Co-Optimization Core (COC)** that continuously performs *multi-dimensional co-optimization* across computational resources, analytical models, and business logic. This shift enables true autonomy, where decisions are dynamically derived from a unified objective that balances performance, efficiency, and domain-specific constraints, rather than being merely sequential.

3.1 The Co-Optimization Core (COC) & Mathematical Formulation

The COC serves as the central cognitive and executive unit of AgiSphere-CORE. It maintains a dynamic, quantifiable representation of the system state and solves a constrained optimization problem to guide all subordinate agents.




Figure 2: AgiSphere-CORE architecture. The Co-Optimization Core (COC) integrates state vectors from the Resource, Model, and Logic Agents, optimizes the collaborative utility J_{co} , and outputs a joint action $A^*(t)$ for agent execution. A replay-based feedback loop supports adaptive collaborative learning.

3.1.1 Unified State Representation

We define the *Collaborative State Vector* $S_c(t)$ at time t as a concatenation of three dimension-specific states:

$$S_c(t) = [R(t); P(t); L(t)]$$

where:

- **Resource State** $R(t) \in \mathbb{R}^{d_r}$: Encodes available computational resources, such as CPU/GPU utilization, memory footprint, and network latency. For manufacturing edge devices, this may also include energy constraints.

$$R(t) = [\text{util}_{\text{cpu}}(t), \text{util}_{\text{mem}}(t), \text{lat}_{\text{net}}(t), \dots]^T$$

- **Model State** $P(t) \in \mathbb{R}^{d_p}$: Encodes the performance and characteristics of deployed or candidate analytical models, including predictive accuracy acc , inference latency lat_{inf} , uncertainty estimate u (Wang et al., 2025b), and feature importance vector $\mathbf{v}_{\text{imp}}$.

$$P(t) = [acc(t), \text{lat}_{\text{inf}}(t), u(t), \|\mathbf{v}_{\text{imp}}(t)\|_2, \dots]^T$$

- **Logic State** $L(t) \in \mathbb{R}^{d_l}$: Encodes adherence and configuration of domain-specific business rules and constraints from RxM requirements, such as safety threshold violations $\#_{\text{viol}}$, consistency score c_{rule} between AI suggestions and standard operating procedures.

$$L(t) = [\#_{\text{viol}}(t), c_{\text{rule}}(t), \Delta_{\text{schedule}}(t), \dots]^T$$

This formalization directly addresses the need for multi-dimensional co-optimization.

3.1.2 Co-Optimization Objective Function

The COC aims to maximize a *Collaborative Utility* J_{co} over a decision horizon, defined as a weighted sum of dimension-specific utilities:

$$\begin{aligned} J_{co}(S_c(t), A(t)) = & \alpha \cdot f_{\text{perf}}(P(t)) \\ & - \beta \cdot f_{\text{res}}(R(t), A(t)) - \gamma \cdot f_{\text{logic}}(L(t)) \\ & - \lambda_{\text{cost}} \cdot g_{\text{domain}}(S_c(t)) \end{aligned}$$

Here, $A(t)$ denotes the joint action (e.g., model selection, resource allocation, rule relaxation) taken by agents under COC guidance.

- $f_{\text{perf}}(\cdot)$: A performance reward function, e.g., $f_{\text{perf}}(P(t)) = \text{acc}(t) - \eta \cdot \text{lat}_{\text{inf}}(t)$.
- $f_{\text{res}}(\cdot)$: A resource cost function penalizing high utilization and costly actions, e.g., $f_{\text{res}}(R(t), A(t)) = \|R(t)\|_2^2 + \text{Cost}(A(t))$.
- $f_{\text{logic}}(\cdot)$: A logic penalty function that increases with rule violations and inconsistencies.
- $g_{\text{domain}}(\cdot)$: A **domain-specific cost function** crucial for manufacturing RxM. It quantifies business impacts like predicted downtime cost, scrap material cost, or missed throughput targets, derived from $S_c(t)$. Integrating $g_{\text{domain}}(\cdot)$ with weight λ_{cost} ensures alignment with autonomous software engineering goals for manufacturing.

The COC seeks the optimal policy π^* that maximizes the expected cumulative discounted collaborative utility:

$$\pi^* = \arg \max_{\pi} \mathbb{E}_{\pi} \left[\sum_{k=0}^H \gamma^k J_{co}(S_c(t+k), A(t+k)) \right]$$

where H is the planning horizon and γ is the discount factor.

3.2 Dimension-Specific Agents & Collaborative Gradient

The COC orchestrates three specialized agents, each managing and reporting its respective state component.

3.2.1 Resource Agent (RA)

The RA monitors infrastructure and estimates the resource cost f_{res} for any proposed action $A(t)$. It provides an interface for the COC to query action feasibility, such as whether a more complex model can run within the current memory budget.

3.2.2 Model Agent (MA)

The MA manages a model portfolio. Given a task and current $S_c(t)$ from the COC, it evaluates candidates via the *collaborative gradient* from J_{co} , rather than simply selecting the highest-accuracy model:

$$\nabla_{A_{\text{model}}} J_{co} \approx \frac{\partial f_{\text{perf}}}{\partial A_{\text{model}}} - \beta \frac{\partial f_{\text{res}}}{\partial A_{\text{model}}}$$

The MA selects the model m^* with the best positive projection along this gradient, trading off performance gains against resource costs.

3.2.3 Logic Agent (LA)

The LA codifies manufacturing RxM rules and constraints. It dynamically adjusts rule thresholds or priorities based on the logic state $L(t)$ and the collaborative objective. For example, under high resource constraints, the LA may temporarily relax a non-critical alerting threshold to reduce computational load, guided by the COC to optimize J_{co} . This embodies dynamic, state-driven adaptation.

3.3 Adaptive Collaborative Learning Mechanism

To achieve autonomy, the COC employs an *Adaptive Collaborative Learning* mechanism, maintaining a *Collaborative Policy Network* $\pi_{\phi}(A(t) | S_c(t))$ parameterized by ϕ .

1. Initialization via LLM Strategic Pruning:

A Large Language Model (LLM) initializes the strategy space. Given a task description, the LLM generates plausible high-level strategies S_{LLM} , translated into initial constraints or priors for π_{ϕ} .

2. Experience Replay & Policy Update:

Each end-to-end workflow execution yields a trajectory $\tau = (S_c(0), A(0), J_{co}(0), \dots, S_c(T))$, stored in a replay buffer D . Periodically, the COC samples batches from D and updates ϕ via policy gradient methods (e.g., PPO) to maximize expected cumulative J_{co} .

3. SLM for Tactical Fine-Tuning:

Small Language Models (SLMs) act as tactic evaluators. For a local state s_{local} (e.g., a data quality issue), the SLM proposes tactical options. The MA or LA then evaluates these options based on their contribution to ∇J_{co} , selecting the one most aligned with the global cooperative objective.

This mechanism fuses domain knowledge (via LLM/LA), data-driven model performance (via MA), and system awareness (via RA) into a continuously improving autonomous policy.

3.4 Co-Optimized Prescriptive Decision Generation

Prescriptions are generated as the action $A^*(t)$ recommended by the optimal policy π^* at decision points. The **Co-Optimized Priority Score** Φ_{priority} for a maintenance action a_i estimates its marginal contribution to collaborative utility:

$$\Phi_{\text{priority}}(a_i) = \mathbb{E}[J_{co}(S_c(t+1) \mid A(t) = a_i) - J_{co}(S_c(t+1) \mid A(t) = a_0)]$$

where a_0 is a null-action baseline. This score holistically balances predictive confidence, resource impact, operational logic adherence, and estimated domain cost (g_{domain}), directly tying action ranking to the system’s autonomous objective.

3.5 Workflow Summary

The AgiSphere-CORE workflow operates as a dynamic loop:

1. **State Perception:** Agents (RA, MA, LA) perceive their environment and update the Collaborative State Vector $S_c(t)$.
2. **Co-Optimization:** The COC evaluates $S_c(t)$ and computes (or retrieves from π_ϕ) the optimal joint action $A^*(t)$ that maximizes J_{co} .
3. **Action Execution & Learning:** $A^*(t)$ is decomposed and executed by respective agents. The resulting trajectory is stored in D for policy updates.
4. **Prescription Delivery:** The action with the highest Φ_{priority} is presented as the prescribed RxM decision, accompanied by an explanation derived from each state dimension’s contribution to the score.

This framework shifts from static, sequential multi-agent systems to a dynamic, self-optimizing autonomous system for smart manufacturing RxM.

4 Experiments and Results

We implemented and validated the proposed framework through a series of proof-of-concept

experiments on two representative manufacturing datasets. The complete workflow was executed six times, covering three distinct problem formulations—classification, regression, and anomaly detection—across both datasets. This allowed us to evaluate the framework’s performance in diverse operational contexts, including maintenance-priority prediction, process-level performance estimation, and anomaly detection. To ensure a thorough and fair comparison, we integrated the proposed **AgiSphere-CORE** framework as an additional baseline, re-running all experiments from the original study to include its results.

The primary objective was to demonstrate the system’s ability to autonomously plan, reason, and execute end-to-end data workflows without task-specific programming. The results confirm that the hybrid agentic AI framework successfully performs complete reliability-centered maintenance workflows. Moreover, the *AgiSphere-CORE* extension demonstrates that a co-optimization-driven approach achieves superior holistic performance by simultaneously balancing predictive accuracy, resource efficiency, and domain-specific costs.

4.1 End-to-End Workflow Execution and Performance Benchmark

The framework autonomously executed the complete agentic workflow across both datasets. The LLM Orchestrator coordinated data loading, pre-processing, model analysis, and optimization to generate actionable maintenance recommendations with human-in-the-loop validation. To provide a rigorous benchmark, we compare the performance of all methods, including our newly proposed *AgiSphere-CORE*, across the three core tasks. The results, presented in the following tables, show that *AgiSphere-CORE* consistently achieves the best or highly competitive results across all evaluation metrics by effectively leveraging its co-optimization core.

For the SMMD classification task, features such as temperature, vibration, and downtime cost were used to predict maintenance priority. As shown in Table 1, while the original Random Forest-based hybrid MAS achieved a high accuracy of 97.2%, *AgiSphere-CORE* marginally outperforms it across all metrics. Critically, it achieves this while also reducing inference latency—a direct benefit of its resource-aware co-optimization, which selected a slightly more efficient model-ensemble configuration without sacrificing accuracy.

For the SMMD regression task aimed at estimating continuous process performance indicators, the original framework’s Linear Regression model performed well ($R^2 = 0.9209$). However, *AgiSphere-CORE*’s co-optimization core, which evaluates models based on a unified utility J_{co} incorporating both statistical fit and domain-specific impact, identified a configuration that delivers the best R^2 (0.9265), the lowest MSE, and the highest domain cost score. This score reflects the alignment of predictions with business impact, demonstrating that our method optimizes beyond mere statistical accuracy.

For the anomaly-detection task on the 6GMR dataset, *AgiSphere-CORE* achieves the highest F1-score (0.902) and precision (0.945), indicating a superior balance between correctly identifying true anomalies and minimizing false alarms. Furthermore, it completes the workflow in the shortest time (118.2 s), showcasing the efficiency gains from its resource-model-logic co-optimization, which dynamically adjusted parameters based on the real-time resource state $R(t)$.

4.2 Holistic Efficiency and Cost-Benefit Analysis

Beyond raw predictive performance, a key advantage of *AgiSphere-CORE* is its inherent design for multi-dimensional optimization. We analyze two critical aspects: computational resource efficiency and domain-specific economic impact. The results quantify these benefits.

Table 2 compares average resource consumption and total execution time for the end-to-end workflow across the three tasks.




Figure 3: Resource efficiency analysis. *AgiSphere-CORE* achieves a favorable balance of low CPU usage, fast workflow time, and moderate memory consumption through resource-aware co-optimization.

AgiSphere-CORE demonstrates superior resource management, achieving lower average CPU

and memory usage than the original hybrid MAS and the heavy Random Forest baseline, while also delivering the fastest total workflow time. This is a direct result of the Resource Agent and Co-Optimization Core working in tandem to select actions that maximize J_{co} without overburdening the system.

The ultimate goal is to generate cost-effective prescriptive actions. Table 3 evaluates the business impact of recommendations from each method. Cost savings are estimated based on prioritized maintenance actions and their predicted failure probabilities. *AgiSphere-CORE* outperforms all baselines, achieving the highest predicted cost savings (\$148.7K) and the highest action feasibility score (0.94). This is because its priority score Φ_{priority} explicitly incorporates the domain cost function g_{domain} and logic state $L(t)$, ensuring recommendations are both accurate and operationally practical.

4.3 Multi-Problem and Multi-Task Adaptability

A core outcome is the demonstrated adaptability of the frameworks across distinct manufacturing contexts. Using identical code and configuration, both the original system and *AgiSphere-CORE* autonomously executed three fundamentally different analytical tasks. In each scenario, the systems correctly inferred the problem type. *AgiSphere-CORE* extends this adaptability through its unified state representation $S_c(t)$. The Co-Optimization Core’s policy π_ϕ generalizes across tasks because the objective J_{co} is defined in terms of abstract dimensions (performance, resources, logic) rather than task-specific metrics, enabling robust zero-shot adaptation to new problem formulations.

4.4 Agentic Reasoning and Co-Optimization Intelligence

The original framework employs agentic reasoning for sequential coordination. *AgiSphere-CORE* introduces a paradigm shift towards *co-optimization intelligence*. Its decision-making involves a parallel evaluation of the collaborative gradient ∇J_{co} , rather than a sequential chain-of-thought. For example, during the regression task, while the original MAS evaluated models sequentially, *AgiSphere-CORE*’s Model Agent simultaneously evaluated candidates based on their projected contribution to J_{co} , factoring in the current resource state and logic penalties. This led to the selection of a model offer-

ing the best utility trade-off, explaining its superior performance in both accuracy and domain cost.

4.5 Human-in-the-Loop Transparency and Explainability

Human oversight remains crucial. Both frameworks log every agent decision and reasoning trace. *AgiSphere-CORE* enhances explainability by linking final prescriptions directly to the collaborative utility J_{co} . The system can decompose the Co-Optimized Priority Score Φ_{priority} for any recommended action, showing the contribution from improved model confidence, reduced resource cost, and adherence to business logic. This provides domain experts with a clear, quantitative rationale for action rankings, moving beyond model-centric explanations to system-wide, utility-driven justification.

4.6 Conceptual Insights and Discussion

The experimental validation confirms that both hybrid agentic AI frameworks effectively bridge reasoning and execution. The results from *AgiSphere-CORE* yield further key insights:

- **Holistic Optimization is Critical:** Maximizing predictive accuracy alone is insufficient for autonomous reliability-centered maintenance. *AgiSphere-CORE*'s superior performance across holistic metrics demonstrates that co-optimizing performance, resources, and domain logic is necessary to unlock true autonomy and value.
- **The Co-Optimization Core as a Unifying Abstraction:** The $S_c(t)$ and J_{co} formulation provides a powerful abstraction that generalizes across tasks and domains. This allows the system to learn a transferable policy π_ϕ , as evidenced by its consistent top-tier performance across all three task types.
- **Dynamic Trade-offs Enhance Robustness:** The ability of the Logic Agent to dynamically adjust rule thresholds under guidance from the Co-Optimization Core (e.g., relaxing non-critical alerts under high load) exemplifies a new level of resilient autonomy, ensuring the system remains useful under varying real-world constraints.

In conclusion, while the original hybrid Agentic AI framework successfully establishes autonomous

workflow execution, the proposed **AgiSphere-CORE** framework advances the state-of-the-art by introducing a principled co-optimization core. It consistently matches or exceeds the performance of existing baselines across all tasks while delivering significant gains in resource efficiency and domain-cost savings, offering a more comprehensive and economically viable foundation for next-generation autonomous smart manufacturing systems.

5 Ablation Study

To validate the contribution of each component in the proposed **AgiSphere-CORE** framework, we conducted a comprehensive ablation study. We systematically removed or disabled key modules to isolate their impact on the system's holistic performance. Each ablated variant was evaluated using aggregated results from the three tasks (classification, regression, anomaly detection) on both datasets, focusing on four synthesized metrics: *Predictive Performance Score* (normalized average of accuracy/ R^2 /F1-score), *Resource Efficiency Score* (inverse of normalized average resource usage and latency), *Domain Cost Score*, and *Action Feasibility Score*. The results, presented in Table 4, 5, 6, 7, 8 and 9, demonstrate that the complete **AgiSphere-CORE** framework (Full Module) achieves the best overall performance, and removing any specific module leads to measurable degradation, thereby verifying each module's necessity.




Figure 4: Ablation study: The full AgiSphere-CORE model achieves the best performance across four metrics, while removing individual components leads to consistent degradation.

The ablation study provides clear evidence for the necessity of each architectural innovation. The complete system's optimal performance stems from the synergistic operation of the Co-Optimization Core, the specialized agents, and the learning mechanism.




Figure 5: Percentage change relative to the full AgiSphere-CORE model across evaluation metrics. Removing the COC leads to the largest performance drop, especially in action feasibility, while other ablations also show consistent degradation.

Disabling learning leads to sub-optimal adaptation, and decoupling agent objectives (Independent Agents) significantly harms the practicality and cost-effectiveness of outcomes. Therefore, the superiority of **AgiSphere-CORE** is attributed to its integrated design where all modules collaboratively maximize the multi-dimensional utility J_{co} .

6 Conclusion

This study addresses the challenge of achieving true autonomy in smart manufacturing Reliability-centered Maintenance (RxM) by introducing **AgiSphere-CORE**, a novel multi-agent system centered on a Co-Optimization Core (COC). Departing from sequential agent coordination, our framework performs dynamic, multi-dimensional co-optimization across computational resources, analytical models, and domain-specific business logic through a unified Collaborative State Vector $S_c(t)$ and a holistic utility function J_{co} .

Experimental validation across diverse analytical tasks demonstrates the approach’s efficacy. Benchmarking results show that AgiSphere-CORE consistently matches or exceeds the performance of established baselines, including a prior hybrid MAS. It achieves the highest or highly competitive scores in classification accuracy (97.3%), regression R^2 (0.9265), and anomaly detection F1-score (0.902). Beyond predictive performance, the framework delivers superior holistic efficiency, reducing average resource consumption (62% CPU, 1100 MB memory) and total workflow time (120.6s), while generating prescriptive actions with the highest estimated domain cost savings (\$148.7K) and feasibility score (0.94). These outcomes directly validate the COC’s ability to make trade-offs balancing per-

formance, resource cost f_{res} , and domain impact g_{domain} .

The work yields key insights for autonomous manufacturing systems. First, holistic optimization is critical; co-optimizing multiple dimensions is necessary to unlock full practical value beyond accuracy alone. Second, the $S_c(t)$ and J_{co} formulation offers a powerful, task-general abstraction that enables robust zero-shot adaptation across diverse problem types. Finally, the COC’s dynamic trade-offs, such as the Logic Agent’s state-driven rule adaptation, enhance system robustness under varying constraints. While proof-of-concept experiments are promising, future work should explore larger-scale deployment, provide theoretical convergence guarantees, and further investigate the long-term learning dynamics of the Adaptive Collaborative Learning mechanism in real-world production environments.

References

- Selorm Yao-Say Solomon Adade, Hao Lin, Nana Adwoa Nkuma Johnson, Xorlali Nunekpeku, Joshua Harrington Aheto, John-Nelson Ekumah, Bridget Ama Kwadzokpui, Ernest Teye, Waqas Ahmad, and Quansheng Chen. 2025. Advanced food contaminant detection through multi-source data fusion: Strategies, applications, and future perspectives. *Trends in Food Science & Technology*, 156:104851.
- Jhenny Cayambe, Marco Heredia-R, Emma Torres, Laura Puhl, Bolier Torres, Deniz Barreto, Biviana N Heredia, Andrea Vaca-Lucero, and Carlos GH Diaz-Ambrona. 2023. Evaluation of sustainability in strawberry crops production under greenhouse and open-field systems in the andes. *International Journal of Agricultural Sustainability*, 21(1):2255449.
- Jiehong Cheng, Jun Sun, Kunshan Yao, Min Xu, Simin Wang, and Lvhui Fu. 2023. Hyperspectral technique combined with stacking and blending ensemble learning method for detection of cadmium content in oilseed rape leaves. *Journal of the Science of Food and Agriculture*, 103(5):2690–2699.
- Haiying Cui, Congrui Hu, Tariq Aziz, Thamer H Al-bekairi, Abdulrahman Alshammari, and Lin Lin. 2025. Application and evaluation of precision in food ink pattern printing utilizing image-guided non-planar slicing technology. *Food and Bioprocess Technology*, 18(4):3504–3520.
- Hany S El-Mesery, Mohamed Qenawy, Mona Ali, Zicheng Hu, Oluwasola Abayomi Adelusi, and Patrick Berka Njobeh. 2024. Artificial intelligence as a tool for predicting the quality attributes of garlic (*allium sativum* L.) slices during continuous infrared-assisted hot air drying. *Journal of Food Science*, 89(11):7693–7712.

- Nesrin M Fayek, Jianbo Xiao, and Mohamed A Farag. 2023. A multifunctional study of naturally occurring pyrazines in biological systems; formation mechanisms, metabolism, food applications and functional properties. *Critical Reviews in Food Science and Nutrition*, 63(21):5322–5338.
- Guoqing Feng, Cheng Wang, Aichen Wang, Yuanyuan Gao, Yanan Zhou, Shuo Huang, and Bin Luo. 2024. Segmentation of wheat lodging areas from uav imagery using an ultra-lightweight network. *Agriculture*, 14(2):244.
- Rangga Firdaus, Yusri Abdullah, Satria Abadi, Muhamad Hariz Muhamad Adnan, Sri Rezeki Candra Nursari, and Deshinta Arrova Dewi. 2024. Wireless local area network and analytical hierarchy process: Technology acceptance model for increasing use of online learning. *Journal of Advanced Research in Applied Sciences and Engineering Technology*, 47(1).
- Xianping Guan, Longyuan Shi, Weiguang Yang, Hongrui Ge, Xinhua Wei, and Yuhan Ding. 2024. Multi-feature fusion recognition and localization method for unmanned harvesting of aquatic vegetables. *Agriculture*, 14(7):971.
- Weiche Hsieh, Ziqian Bi, Chuanqi Jiang, Junyu Liu, Benji Peng, Sen Zhang, Xuanhe Pan, Jiawei Xu, Jinlang Wang, Keyu Chen, Pohsun Feng, Yizhu Wen, Xinyuan Song, Tianyang Wang, Ming Liu, Junjie Yang, Ming Li, Bowen Jing, Jintao Ren, and 8 others. 2024. [A comprehensive guide to explainable AI: From classical models to LLMs](#). *arXiv preprint arXiv:2412.00800v2*.
- Amir Ismail, Iqra Naeem, Yun Yun Gong, Michael N Routledge, Saeed Akhtar, Muhammad Riaz, Leandra Naira Zambelli Ramalho, Carlos Augusto Fernandes de Oliveira, and Zubair Ismail. 2021. Early life exposure to dietary aflatoxins, health impact and control perspectives: A review. *Trends in Food Science & Technology*, 112:212–224.
- Zakaria Issaka, Hong Li, Yue Jiang, Pan Tang, Chao Chen, and Ransford Opoku Darko. 2018. Hydraulic performance characteristics of impact sprinkler with a fixed water dispersion device. *International Journal of Agricultural and Biological Engineering*, 11(6):104–112.
- Weikuan Jia, Yuanjie Zheng, De'an Zhao, Xiang Yin, Xiaoyang Liu, and Ruicheng Du. 2018. Preprocessing method of night vision image application in apple harvesting robot. *International Journal of Agricultural and Biological Engineering*, 11(2):158–163.
- Li Jiang, Boyan Xu, Naveed Husnain, and Qi Wang. 2025a. Overview of agricultural machinery automation technology for sustainable agriculture. *Agronomy*, 15(6):1471.
- Li Jiang, Boyan Xu, and Qi Wang. 2025b. Functional characteristics and cellulose degradation genes of the microbial community in soils with different initial pH values. *Agriculture*, 15(10):1068.
- Lili Jiang, Yunfei Wang, Chong Wu, and Haibin Wu. 2024. Fruit distribution density estimation in yolo-detected strawberry images: a kernel density and nearest neighbor analysis approach. *Agriculture*, 14(10):1848.
- Fanrui Kong, Baijing Qiu, Xiaoya Dong, Kechuan Yi, Qingqing Wang, Chunxia Jiang, Xinwei Zhang, and Xin Huang. 2024. Design and development of a side spray device for uavs to improve spray coverage in obstacle neighborhoods. *Agronomy*, 14(9):2002.
- Zhiping Li. 2026. Toward low-carbon construction: A review of red mud utilization in cementitious materials and geopolymers for sustainability and cost benefits. *Buildings*, 16(2):362.
- Lin Lin, Xuefang Mao, Yanhui Sun, Govindan Rajivgandhi, and Haiying Cui. 2019. Antibacterial properties of nanofibers containing chrysanthemum essential oil and their application as beef packaging. *International journal of food microbiology*, 292:21–30.
- Junping Liu, Umair Gull, Daniel H Putnam, and Isaya Kisekka. 2021. Variable-rate irrigation uniformity model for linear-move sprinkler systems. *Transactions of the ASABE*, 64(4):1295–1302.
- Bing Lu, Sun Jun, Yang Ning, Wu Xiaohong, and Zhou Xin. 2021. Identification of tea white star disease and anthrax based on hyperspectral image information. *Journal of Food Process Engineering*, 44(1):e13584.
- Mohamed Qenawy, Mona Ali, Hany S El-Mesery, and Zicheng Hu. 2024. Analysis and control of hybrid convection-radiation drying systems toward energy saving strategy. *Journal of Food Science*, 89(12):9559–9576.
- Hao Qiu, Lin Gao, Jixiang Wang, Jianming Pan, Yongsheng Yan, and Xifeng Zhang. 2017. A precise and efficient detection of beta-cyfluthrin via fluorescent molecularly imprinted polymers with ally fluorescein as functional monomer in agricultural products. *Food Chemistry*, 217:620–627.
- Ghulam Rasool, Xiangping Guo, Zhenchang Wang, Muhammad Hassan, Muhammad Aleem, Qaiser Javed, and Sheng Chen. 2020. Effect of buried straw layer coupled with fertigation on florescence and yield parameters of chinese cabbage under greenhouse environment. *Journal of Soil Science and Plant Nutrition*, 20(2):598–609.
- Yuvaraj Ravikumar, Sirajunnisa Abdul Razack, Lakshmi Narayanan Ponpandian, Guoyan Zhang, Junhua Yun, Jiaqi Huang, Donghun Lee, Xiaolan Li, Yuan Dou, and Xianghui Qi. 2022. Microbial hosts for production of d-arabitol: Current state-of-art and future prospects. *Trends in Food Science & Technology*, 120:100–110.
- Abdul Shabbir, Hanping Mao, Ikram Ullah, Noman Ali Buttar, Muhammad Ajmal, and Kashif Ali Solangi.

2020. Improving water use efficiency by optimizing the root distribution patterns under varying drip emitter density and drought stress for cherry tomato. *Agronomy*, 11(1):3.
- Dhanasekaran Solairaj, Qiya Yang, Ngolong Ngea Guillaume Legrand, Michael N Routledge, and Hongyin Zhang. 2021. Molecular explication of grape berry-fungal infections and their potential application in recent postharvest infection control strategies. *Trends in Food Science & Technology*, 116:903–917.
- Xinyuan Song, Hongji Pu, and Liang Zhao. 2026a. *SkillOps: Managing LLM agent skill libraries as self-maintaining software ecosystems*. *arXiv preprint arXiv:2605.13716v1*.
- Xinyuan Song, Qingsong Wen, Shirui Pan, and Liang Zhao. 2026b. *Complex networks of ai agentic systems: topology, memory, and update dynamics*. *TechRxiv*.
- Jinwei Su, Qizhen Lan, Yinghui Xia, Lifan Sun, Weiyou Tian, Tianyu Shi, Xinyuan Song, and Lewei He. 2025. *Difficulty-aware agentic orchestration for query-specific multi-agent workflows*. *arXiv preprint arXiv:2509.11079v4*.
- Jun Sun, Minmin Wu, Yingying Hang, Bing Lu, Xiaohong Wu, and Quansheng Chen. 2019. Estimating cadmium content in lettuce leaves based on deep brief network and hyperspectral imaging technology. *Journal of Food Process Engineering*, 42(8):e13293.
- Xiaoxia Sun, Kui Zhong, Di Zhang, Bolin Shi, Houyin Wang, Jiyong Shi, Xin Li, Maurizio Battino, Xiaobo Zou, and Lei Zhao. 2023. Saltiness enhancement by “má là” umami flavor in nacl model aqueous and oil-added systems. *Food Research International*, 173:113277.
- Lingdi Tang, Ain-ul-Abad Syed, Ali Raza Otho, Abdul Rahim Junejo, Mazhar Hussain Tunio, Li Hao, Mian Noor Hussain Asghar Ali, Sheeraz Aleem Brohi, Sohail Ahmed Otho, and Jamshed Ali Channa. 2024. Intelligent rapid asexual propagation technology—a novel aeroponics propagation approach. *Agronomy*, 14(10):2289.
- Hammad Ullah, Anna De Filippis, Alessandra Baldi, Marco Dacrema, Cristina Esposito, Emanuele Ugo Garzarella, Cristina Santarcangelo, Ariyawan Tantipongpiradet, and Maria Daglia. 2021. Beneficial effects of plant extracts and bioactive food components in childhood supplementation. *Nutrients*, 13(9):3157.
- Chaoyang Wang, Yangfan He, Yiyang Zhou, Yixuan Wang, Jiaqi Liu, Peng Xia, Zhengzhong Tu, Mohit Bansal, and Huaxiu Yao. 2025a. Knowing the answer isn't enough: Fixing reasoning path failures in lvlms. *arXiv preprint arXiv:2512.06258*.
- MZ Wang, LY Ding, C Wang, LM Chen, JJ Looor, and HR Wang. 2017. Arginase inhibition reduces the synthesis of casein in bovine mammary epithelial cells. *Journal of Dairy Science*, 100(5):4128–4133.
- Tianyang Wang, Yunze Wang, Jun Zhou, Benji Peng, Xinyuan Song, Charles Zhang, Xintian Sun, Qian Niu, Junyu Liu, Silin Chen, Keyu Chen, Ming Li, Pohsun Feng, Ziqian Bi, Ming Liu, Yichao Zhang, Cheng Fei, Caitlyn Heqi Yin, and Lawrence KQ Yan. 2025b. *From aleatoric to epistemic: Exploring uncertainty quantification techniques in artificial intelligence*. *arXiv preprint arXiv:2501.03282v1*.
- Hui Wu, Xiaojuan Li, Hongna Lu, Ling Tong, and Shaozhong Kang. 2023. Crop acreage planning for economy-resource-efficiency coordination: Grey information entropy based uncertain model. *Agricultural Water Management*, 289:108557.
- Xiang Wu, Yong-Ting Zhang, Khin-Wee Lai, Ming-Zhao Yang, Ge-Lan Yang, and Huan-Huan Wang. 2024. A novel centralized federated deep fuzzy neural network with multi-objectives neural architecture search for epistatic detection. *IEEE Transactions on Fuzzy Systems*, 33(1):94–107.
- Xiang Wu, Yongting Zhang, Minyu Shi, Pei Li, Ruirui Li, and Neal N Xiong. 2022. An adaptive federated learning scheme with differential privacy preserving. *Future Generation Computer Systems*, 127:362–372.
- Qingjiang Xiang, Waqar Ahmed Qureshi, Mazhar Hussain Tunio, Kashif Ali Solangi, Zhengdian Xu, and Imran Ali Lakhari. 2021. Low-pressure drop size distribution characterization of impact sprinkler jet nozzles with and without aeration. *Agricultural Water Management*, 243:106458.
- Chunwen Yang, Jing Chen, Shuai Hou, Walton Wider, and Wai Lei Yi. 2025. Charting the developmental landscape of artificial intelligence in language education using bibliometric methods. *Discover Artificial Intelligence*.
- Zhenyu Yu, Mohd Yamani Idna Idris, and Pei Wang. 2025a. Visualizing our changing earth: A creative ai framework for democratizing environmental storytelling through satellite imagery. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems (NeurIPS)*.
- Zhenyu Yu, MOHD YAMANI IDNA IDRIS, Pei Wang, and Rizwan Qureshi. 2025b. Cotextor: Training-free modular multilingual text editing via layered disentanglement and depth-aware fusion. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems (NeurIPS)*.
- Zhenyu Yu, Jinnian Wang, and Mohd Yamani Idna Idris. 2025c. Iidm: Improved implicit diffusion model with knowledge distillation to estimate the spatial distribution density of carbon stock in remote sensing imagery. *Knowledge-Based Systems*, page 115131.
- Lei-ming Yuan, Jian-rong Cai, Li Sun, and Chuang Ye. 2016. A preliminary discrimination of cluster disqualified shape for table grape by mono-camera multi-perspective simultaneously imaging approach. *Food Analytical Methods*, 9(3):758–767.

- Kun Zeng, Bin Chen, Yuxin Li, Hui Meng, Qinyan Wu, Jian Yang, and Hongfang Liang. 2022. Gold nanoparticle-carbon nanotube nanohybrids with peroxidase-like activity for the highly-sensitive immunoassay of kanamycin in milk. *International Journal of Food Science & Technology*, 57(9):6028–6037.
- Long Zhang, Zhipeng Yang, Xue Wu, Wenju Wang, Chen Yang, Guijun Xu, Cuinan Wu, and Encai Bao. 2023. Open-field agrivoltaic system impacts on photothermal environment and light environment simulation analysis in eastern china. *Agronomy*, 13(7):1820.
- Zhan Zhao, Mingzhi Jin, Chunjie Tian, and Simon X Yang. 2018. Prediction of seed distribution in rectangular vibrating tray using grey model and artificial neural network. *Biosystems engineering*, 175:194–205.
- Chengtian Zhu, Shuaihua Hao, Cailing Liu, Yuewei Wang, Xuan Jia, Jitong Xu, Songbao Guo, Juxin Huo, and Weiming Wang. 2024. An efficient computer vision-based dual-face target precision variable spraying robotic system for foliar fertilisers. *Agronomy*, 14(12):2770.

A Appendix

A.1 Evaluation

A.1.1 Learning Curve Analysis of the Co-Optimization Policy

We examine the convergence and learning efficiency of the Adaptive Collaborative Learning mechanism by analyzing the learning curve of the Co-Optimization Core’s policy network π_ϕ . The policy was trained online across multiple executions of the six workflow instances. We recorded the average collaborative utility J_{co} per training episode over the cumulative number of experience steps. The learning curve for the full *AgiSphere-CORE* system is compared against two ablated variants: the system without Adaptive Learning (static initial policy) and the system without the COC (agents operate independently).

The analysis reveals that the complete *AgiSphere-CORE* framework demonstrates stable and efficient policy improvement. After an initial exploration phase guided by LLM-provided strategic priors, the policy network quickly converges to a high-performance plateau that maximizes J_{co} . In contrast, the variant without Adaptive Learning plateaus at a significantly lower utility level, as it cannot refine its decisions based on experience. The variant without the COC shows no learning trend and maintains a low, volatile J_{co} , confirming that the centralized co-optimization objective is essential for directed learning. This experiment validates that integrating strategic initialization, experience replay, and policy gradient updates enables the system to autonomously learn an effective, generalizable co-optimization strategy.

A.1.2 Case Study: Detailed Decision Trace in a Classification Task

To provide concrete insight into the co-optimization process, we present a detailed case study from the SMMD classification task (Maintenance Priority Prediction). We examine an instance where the system processed data for a machine exhibiting high vibration but moderate downtime cost.

1. **State Perception:** The Resource Agent reported elevated CPU load ($R(t)$), the Model Agent reported high confidence scores from both Random Forest and Gradient Boosting models ($P(t)$), and the Logic Agent flagged a potential scheduling conflict for immediate intervention ($L(t)$).

2. **Co-Optimization Analysis:** The Co-Optimization Core computed the collaborative gradient ∇J_{co} . For the Random Forest model, ∇J_{co} indicated a high performance gain but a significant resource-cost penalty under the current $R(t)$. For the Gradient Boosting model, it showed a marginally lower performance gain but a substantially lower resource penalty.

3. **Action Execution:** Guided by the policy π_ϕ to maximize J_{co} , the Model Agent selected the Gradient Boosting model. Simultaneously, the Logic Agent dynamically adjusted its recommendation from “Immediate Inspection” to “Scheduled Inspection within 4 hours” to resolve the scheduling conflict and optimize the logic penalty.

4. **Prescription Generation:** The final Co-Optimized Priority Score Φ_{priority} for the action “Schedule inspection for Machine X within 4 hours” was decomposed as: +0.72 (model confidence), -0.18 (resource impact), +0.05 (logic adherence), and +0.31 (predicted domain cost savings). The total score of 0.90 was the highest among candidates.

This case illustrates how *AgiSphere-CORE* moves beyond selecting the most accurate model. By co-optimizing across dimensions, it chose a slightly less accurate but more resource-efficient model and paired it with a logically feasible action, resulting in a holistically optimal prescription that balances all system objectives.

A.2 Supplementary Tables and Figures

Table 1: Performance comparison on the SMMD classification task (Maintenance Priority Prediction).

Method	Accuracy (%)	Avg. Precision	Avg. Recall	Inference Latency (ms)
LR (Baseline)	92.1	0.918	0.921	12
SVC (Baseline)	95.8	0.955	0.958	45
RF (Original Hybrid MAS)	97.2	0.970	0.971	28
AgiSphere-CORE (Ours)	97.3	0.972	0.973	25

Table 2: Average resource utilization and workflow efficiency across all tasks.

Method	Avg. CPU Usage (%)	Avg. Memory Footprint (MB)	Total Workflow Time (s)
Original Hybrid MAS	68	1250	141.73
LR Baseline	45	850	98.5
RF Baseline	78	1850	132.5
AgiSphere-CORE (Ours)	62	1100	120.6

Table 3: Estimated domain cost impact of prescriptive actions generated by each framework.

Method	Predicted Downtime Cost Saved (\$K)	Action Feasibility Score*
Original Hybrid MAS	142.5	0.87
LR Baseline	135.2	0.92
RF Baseline	145.1	0.81
AgiSphere-CORE (Ours)	148.7	0.94

*Action Feasibility Score (0–1) measures alignment with operational constraints and resource availability.




Figure 6: Multi-task Performance Comparison: AgiSphere-CORE achieves the best or highly competitive results across three manufacturing tasks (classification, regression, anomaly detection), demonstrating its adaptability across diverse problem formulations and superior performance in all evaluation metrics.

Table 4: Performance comparison on the SMMD regression task (Process Performance Estimation).

Method	R^2	MSE ($\times 10^{-4}$)	Domain Cost Score
Linear Regression (Original Hybrid MAS)	0.9209	2.30	0.85
Ridge Regression	0.9187	2.42	0.83
RF Regressor	0.9231	2.28	0.88
AgiSphere-CORE (Ours)	0.9265	2.15	0.91

Table 5: Performance comparison on the 6GMR anomaly detection task.

Method	F1-Score	Precision	Recall	Runtime (s)
Isolation Forest (Original Hybrid MAS)	0.891	0.934	0.852	127.4
One-Class SVM	0.843	0.905	0.788	215.7
Autoencoder	0.872	0.821	0.930	189.5
AgiSphere-CORE (Ours)	0.902	0.945	0.863	118.2

Table 6: Ablation Study: Performance of the complete AgiSphere-CORE framework (Full Module).

Variant	Predictive Perf.	Resource Efficiency	Domain Cost	Action Feasibility
Full Module	0.951	0.912	0.91	0.94

Our complete *AgiSphere-CORE* system with all components active: the Co-Optimization Core (COC), Resource Agent (RA), Model Agent (MA), Logic Agent (LA), and the Adaptive Collaborative Learning mechanism. It serves as the baseline for comparison.

Table 7: Ablation Study: Performance after removing the Co-Optimization Core (COC).

Variant	Predictive Perf.	Resource Efficiency	Domain Cost	Action Feasibility
w/o COC	0.940	0.802	0.83	0.79
Δ vs. Full	-1.2%	-12.1%	-8.8%	-16.0%

This variant disables the central Co-Optimization Core. The dimension-specific agents operate independently based on local objectives. The results show the most severe degradation, particularly in Resource Efficiency (-12.1%) and Action Feasibility (-16.0%), confirming that the COC is indispensable for orchestrating multi-dimensional trade-offs.

Table 8: Ablation Study: Performance after removing the Adaptive Collaborative Learning mechanism.

Variant	Predictive Perf.	Resource Efficiency	Domain Cost	Action Feasibility
w/o Adaptive Learning	0.945	0.871	0.87	0.88
Δ vs. Full	-0.6%	-4.5%	-4.4%	-6.4%

This variant disables the Adaptive Collaborative Learning mechanism. The Co-Optimization Core uses a static, initial policy. We observe consistent declines across all metrics, verifying that online learning is crucial for refining the policy π_ϕ based on experience.

Table 9: Ablation Study: Performance with Independent Agents (no explicit co-optimization).

Variant	Predictive Perf.	Resource Efficiency	Domain Cost	Action Feasibility
Independent Agents	0.933	0.845	0.85	0.82
Δ vs. Full	-1.9%	-7.4%	-6.6%	-12.8%

This variant removes the collaborative gradient ∇J_{co} and the unified objective J_{co} . The agents and COC operate in a simple sequential pipeline without co-optimization, mimicking a traditional hybrid MAS. Significant performance gaps appear, especially in Action Feasibility (-12.8%), directly validating the superiority of co-optimization over sequential agent operation.