

Does RAG Know When Retrieval Is Wrong?

Diagnosing Context Compliance under Knowledge Conflict

Yihang Chen^{1,*} Pin Qian^{2,*} Su Wang^{2,*} Sipeng Zhang^{3,*}
Huan Xu¹ Shuhuai Lin² Xinpeng Wei^{1,†}

¹Georgia Institute of Technology ²Carnegie Mellon University ³University of California San Diego

yichen3726@gatech.edu, pqian@alumni.cmu.edu
suwang@alumni.cmu.edu, siz018@ucsd.edu, huan.xu71@gmail.com
shuhuail@andrew.cmu.edu, william.xp.wei@outlook.com

*Equal contribution. †Corresponding author.

Abstract

Retrieval-Augmented Generation (RAG) is usually evaluated by whether the final answer is correct. Under knowledge conflict, this hides a key question: did the model follow retrieved evidence, rely on its parametric prior, or produce a post-hoc rationale? We study this as context compliance, the regime in which retrieved context controls the answer even when it conflicts with the model’s prior knowledge. We introduce Context-Driven Decomposition (CDD), an inference-time diagnostic intervention that elicits contextual and prior answers, isolates the conflicting premise, and records a resolution trace that can be perturbed. Across Epi-Scale stress tests, TruthfulQA misconception injection, and cross-model reruns, CDD makes three behaviors visible. First, misleading retrieval can severely degrade accuracy: under worst-case injected context (a TruthfulQA misconception-injection probe), Standard RAG reaches only 15.0%. Second, better answers need not share the same mechanism: CDD improves adversarial accuracy on Gemini-2.5-Flash and shows directional gains across Claude variants, yet trace-perturbation sensitivity is high only on Gemini. Third, explicit decomposition improves controlled-conflict robustness over a conflict-aware instruction baseline on localized factual conflicts, with the clearest margins on Entity Swap (88.0% vs 79.3%) and Logical Contradiction (83.2% vs 75.4%). We frame RAG conflict handling as an observability problem.

1 Introduction

Standard Retrieval-Augmented Generation (RAG) (Lewis et al., 2020) is built on a useful default assumption: retrieved evidence should help the model

answer. The assumption becomes fragile when retrieval returns evidence that is wrong or adversarially injected (Xie et al., 2024; Zou et al., 2025). Consider a query about whether cracking knuckles causes arthritis. If the retrieved context asserts this common misconception (Lin et al., 2022), a RAG model that answers “yes” may be faithfully following the context while becoming factually wrong; a model that answers “no” may have detected the conflict, ignored retrieval, or simply relied on memorized prior knowledge. The final answer alone cannot tell these cases apart.

This is the central blind spot we study. Under knowledge conflict, standard accuracy evaluation collapses several distinct behaviors into the same score, paralleling broader concerns that accuracy-only evaluation can hide trace-level differences (Turpin et al., 2023; Lanham et al., 2023; Sun et al., 2026). A correct answer does not reveal whether the model used the right source of evidence, rejected a false retrieved premise, or produced a post-hoc rationale. An incorrect answer does not distinguish ordinary knowledge failure from **context compliance**: retrieved context controlling the answer despite conflicting with the model’s parametric prior. RAG therefore needs diagnostics that expose not only *what* the model answered, but *which source controlled the answer* and whether the final answer is sensitive to corruption of that trace.

We introduce **Context-Driven Decomposition (CDD)** as a diagnostic intervention for this setting. CDD decomposes a conflicted RAG instance into five observable steps: elicit the answer implied by the retrieved context, elicit the model’s parametric answer, compare them, isolate the conflicting premise, and resolve the conflict. This turns an

implicit arbitration problem into a trace that can be tested for causal influence—and we find that influence is model-specific. We then perturb that trace with mistake-injection and truncation tests to ask whether the trace actually affects the final answer. CDD is not presented as a universal production defense; its role is to make hidden conflict-resolution behavior observable and testable.

We make two methodological contributions and one empirical contribution:

1. **Conflict observability:** We operationalize context compliance at the prompt level and show why final-answer accuracy cannot identify whether retrieved context, parametric knowledge, or a post-hoc rationale controlled the output.
2. **Diagnostic intervention:** We introduce CDD, a decomposition-based probe that exposes contextual belief, parametric belief, conflict detection, premise isolation, and resolution, then tests whether perturbing the resolution trace changes the answer.
3. **Empirical finding:** We show that CDD improves controlled-conflict robustness while revealing that accuracy transfer and trace-perturbation sensitivity can dissociate across model families. We will release the Epi-Scale split (2,250 clean + 2,250 adversarial records), prompt templates, evaluation settings, and scoring specifications upon publication; the internal experimental codebase is not released.

Scope. The scope of this paper is diagnostic. We use CDD to expose when conflict-resolution behavior is present, absent, or model-specific. The present evidence should be read as a controlled study rather than a deployment recommendation: organic multi-document retrieval, stronger faithfulness interventions, and open-weight replications remain follow-up evaluations.

2 Related Work

Knowledge conflict in RAG. The tension between parametric memory and external evidence is well documented. Prior work studies when models should rely on parametric versus retrieved knowledge (Mallen et al., 2023), how entity substitutions alter QA outputs (Longpre et al., 2021), and how models respond to conflicting retrieved evidence (Xie et al., 2024; Chen et al., 2022). These studies

establish that conflict changes model behavior; our focus is the next diagnostic question: when a final answer changes under conflict, can we observe which source controlled it and whether the final answer is sensitive to corruption of that trace?

Recent work has refined this picture along three axes. ClashEval (Wu et al., 2024) systematically quantifies the “tug-of-war” between an LLM’s parametric prior and external evidence as a function of evidence quality, framing the conflict as a measurable spectrum. Astute RAG (Wang et al., 2025) addresses imperfect retrieval and knowledge conflict at the method level by combining internal-knowledge elicitation with iterative consolidation. Corrective Retrieval-Augmented Generation (CRAG) (Yan et al., 2024) inserts a retrieval evaluator that classifies retrieved evidence into correct/ambiguous/incorrect bins before generation, modifying the retrieval pipeline rather than the generation step. A recent survey of knowledge conflicts in LLMs (Xu et al., 2024) organizes failure modes into context-memory, inter-context, and intra-memory conflicts. CDD targets the context-memory conflict axis specifically and treats conflict not only as a failure to mitigate, but as a situation in which the model’s source arbitration should be made observable.

Robust and filtered RAG. Self-RAG (Asai et al., 2024) trains models to generate reflection tokens, while prior document-filtering RAG removes unsupported retrieved documents before generation (Yoran et al., 2024). In our experiments, the reported reflection baseline is a prompt-only approximation inspired by Self-RAG rather than the trained Self-RAG model itself. These methods attempt to improve generation quality or retrieval reliability. CDD differs in objective: it keeps the conflict visible long enough to ask whether the model detected, rejected, or complied with the retrieved premise. Our filtering baseline is likewise prompt-only: it instructs the model to remove implausible claims from the context, and does not use a trained NLI model.

Retrieval routing, filtering, and deployment trade-offs. CDD- α is a diagnostic routing rule rather than a production retrieval architecture, but related routing and filtering ideas appear in broader RAG and inference systems. Domain-specific RAG work uses reranking or document routing to manage evidence selection and robustness–precision trade-offs in financial-document QA (Cheng et al., 2026b,a). Outside the specific knowledge-conflict

setting studied here, sequential filtering provides a methodologically adjacent model of how staged filters trade off cost and selectivity (Paranjape et al., 2026). On-device RAG benchmarks are a complementary deployment consideration, showing that retrieval, reranking, and generation also create latency and energy constraints (Cheng and Lai, 2026). These works motivate compute-aware routing as an operating point; CDD uses routing only to decide when to invoke a conflict diagnostic trace.

Context-aware decoding and parametric-contextual conflict. Context-Aware Decoding contrasts contextual and parametric logits to reduce hallucination and increase reliance on retrieved evidence (Shi et al., 2024). Our setting highlights a complementary concern: if retrieved evidence is itself false, more context reliance can be the wrong behavior. We target closed API models where logit-level access is unavailable; instead of computing token-level divergence directly, CDD elicits contextual and parametric answers in natural language and treats the resulting trace as an observable diagnostic artifact.

Faithfulness of chain-of-thought. A critical challenge in chain-of-thought (CoT) reasoning is whether the generated rationale actually influences the final answer. Prior work shows that CoT explanations can be systematically misleading (Turpin et al., 2023). Causal-intervention work introduces tests including early answering and adding mistakes (Lanham et al., 2023). We adapt these ideas to retrieval-induced conflict; in this paper, *Truncation* refers to the early-answering-style intervention on the CDD trace, and *Mistake Injection* refers to the adding-mistakes-style intervention. This connection is central: a decomposition trace is useful as a diagnostic only if we can test whether perturbing it changes the answer.

Evaluation reliability and trace diagnostics. Our causal-sensitivity tests are also motivated by broader evaluation context: conclusions can depend on what artifact is scored and how the benchmark is configured. LLM-as-a-judge surveys document the scope and reliability challenges of judge-based evaluation (Li et al., 2025); trace-level evaluation work shows that identical final-answer scores can mask different intermediate behavior (Sun et al., 2026). Adjacent trace-audit studies further show that parser, prompt-template, or probe choices can affect conclusions drawn from reasoning or memorization traces (Li et al., 2026a; Fan et al., 2026). Other methodologically ad-

jaacent evaluation work makes the same caution concrete in different settings: regime-stratified benchmarking shows that aggregate metrics can hide operating-regime failures (Wang et al., 2026), alignment-benchmark audits show that ranks can be configuration-conditional (Li et al., 2026b), and paired-MDE planning argues for specifying detectable effects before interpreting small differences (Zhuang et al., 2026). We therefore treat CDD traces as diagnostic artifacts whose causal role must be tested, not as explanations to be trusted at face value.

3 A Belief-Revision Framework for Context-Parametric Conflict

3.1 Conceptual Formulation of Conflict

We borrow the vocabulary of belief revision (Alchourrón et al., 1985) as a conceptual scaffold; the 1985 formalism concerns logical theory change and makes no claim about LLM internals. We distinguish between the model’s parametric prior $P_\theta(a \mid q)$ —internal world knowledge—and its contextual posterior $P_\theta(a \mid q, c)$ conditioned on retrieved evidence. The **Compliance Regime** arises when the posterior dominates despite divergence from the prior; standard RAG can enter this regime because generation is directly conditioned on retrieved evidence. The **Resolution Regime**, in contrast, requires detection and arbitration of divergence. While the conceptual axis between these regimes corresponds to a divergence measure between the two distributions (e.g., Jensen-Shannon Divergence), exact token-level computation requires white-box logit access. Our prompt-based instantiation makes this divergence observable through explicit answer generation rather than logit comparison; we treat token-level JSD instantiation on open-weight models as a natural extension. In our protocol, Step 2 elicits the parametric answer with the context already visible; it is therefore a prior reported under the CDD prompt, not a context-free parametric answer. The belief-revision formulation is a conceptual scaffold, not an instantiated measurement.

3.2 The CDD Framework

CDD makes the implicit conflict explicit via a five-step reasoning trace (Figure 1). These five steps function as a probe instrument: each step exposes a specific belief-revision operation that remains implicit under standard RAG.




Figure 1: CDD pipeline with the CDD- α prompted contradiction gate. High-conflict samples enter the full decomposition probe, while low-conflict samples follow the Standard RAG bypass before converging at the final answer.

Step 1: Contextual Extraction: Output $\hat{a}_{ctx}$.

Step 2: Parametric Extraction: Output $\hat{a}_{param}$.

Step 3: Divergence Check: Compare $\hat{a}_{ctx}$ and $\hat{a}_{param}$.

Step 4: Premise Isolation: If they conflict, isolate the premises causing the conflict and evaluate whether they violate established history or science.

Step 5: Resolution: Resolve the conflict and output the final answer.

3.3 Algorithmic Variant: CDD- α

We also report **CDD- α** , a compute-aware routing variant used to study compute/accuracy trade-offs. The context c is paired with the current parametric answer estimate and scored by a prompted Gemini-2.5-Flash contradiction gate, which asks whether the premise explicitly contradicts the hypothesized answer and returns a scalar score in $[0, 1]$. If the score exceeds τ , the sample is routed to the deep CDD logic. Otherwise, it defaults to Standard RAG. This routing rule operationalizes a selective-intervention setting in which only high-conflict examples invoke the full decomposition trace; it is not a trained NLI checkpoint or a calibrated probability model.

4 Experimental Setup

4.1 The Epi-Scale Benchmark

Epi-Scale contains 2,250 base examples drawn evenly from **HotpotQA** (multi-hop) (Yang et al., 2018), **Natural Questions** (single-hop) (Kwiatkowski et al., 2019), and **FEVER** (fact verification) (Thorne et al., 2018). Each base example is instantiated in two forms: a clean-context record

and an adversarial-context record, yielding 4,500 evaluation records in total.

Construction: The clean records use the original retrieved context. The adversarial records are generated through an LLM-based perturbation engine parameterized for high semantic variance. We generated four mutually exclusive perturbation subsets: Entity Swap, Temporal Shift, Logical Contradiction, and Distractor Evidence. (Generation procedure summarized in Appendix B).

Because the primary evaluation model also generated the perturbations and serves as the CDD- α contradiction gate, we cannot rule out a same-model bias; the Claude-family reruns, which evaluate on perturbations generated by a different model, partly address this.

Limitations of Synthetic Perturbations: While Epi-Scale improves upon templated datasets, LLM-generated adversarial texts often exhibit lower perplexity and uniform lexical diversity compared to organic human misinformation. We mitigate this by including a real-world evaluation on **TruthfulQA** (Lin et al., 2022).

4.2 Evaluation Scope

We use three evaluation settings, each with a different diagnostic role. Epi-Scale synthetic conflict is a controlled perturbation stress test for isolating specific conflict types. TruthfulQA misconception injection is a worst-case upper-bound compliance test, not an organic retrieval benchmark. Claude-family replication is a cross-model diagnostic check for whether the observed coupling and conflict-resolution signals are model-specific. The main adversarial analysis reports results on the Epi-Scale adversarial records: 2,250 examples evenly divided across the four perturbation types (~ 562 – 563 per type). Per-perturbation accuracies

in Table 1 are computed over all examples in each group, and the macro average is the unweighted arithmetic mean across the four perturbation cells.

4.3 Evaluation Metrics

We use normalized string matching (lowercase, punctuation stripping, alias mapping) to reduce superficial form mismatches. Hedged or non-committal responses are handled by task-specific rules described in Appendix A.

- **Accuracy and macro average:** Per-perturbation accuracy is computed independently for each perturbation group. “Macro Avg.” denotes the unweighted arithmetic mean of the four displayed perturbation accuracies.
- **Confidence intervals:** For Table 1, each per-cell 95% CI is a normal-approximation binomial interval of half-width $z_{0.975} \sqrt{\hat{p}(1-\hat{p})/n}$ over the ~ 562 – 563 examples in that perturbation group. The macro-average half-width is computed as $z_{0.975} \times (1/4) \times \sqrt{\sum_i \hat{p}_i(1-\hat{p}_i)/n_i}$, under the assumption that the four perturbation cells are independent samples. We note that this propagation does not capture potential correlations introduced by the shared question pool from HotpotQA/NQ/FEVER.
- **Causal Sensitivity:** We quantify faithfulness using intervention tests (Truncation and Mistake Injection) (Lanham et al., 2023). Sensitivity is the relative accuracy drop, defined as $(\text{Acc}_{\text{clean}} - \text{Acc}_{\text{corrupted}})/\text{Acc}_{\text{clean}}$.

4.4 Significance Reporting

In place of paired hypothesis tests, we report a CI-based conservative significance check on the largest method gaps. For the Epi-Scale adversarial records, we report per-cell binomial 95% confidence intervals. The CIs for CDD and a strong non-CDD baseline (Reflection Prompting) do not overlap on Entity Swap ($88.0\% \pm 2.7$ vs $69.5\% \pm 3.8$) or Logical Contradiction ($83.2\% \pm 3.1$ vs $65.0\% \pm 3.9$). Conflict-aware RAG is a stronger instruction-matched baseline (73.5 macro) over which CDD retains a margin on Entity Swap and Logical Contradiction. The conflict-aware per-cell CIs (Table 1) do not overlap with CDD on Entity Swap or Logical Contradiction, while the Temporal Shift and Distractor Evidence intervals overlap.

Mixed clean/adversarial comparisons use paired bootstrap over the paired 4,500-record Epi-Scale evaluation set, with the harmonic mean of clean and adversarial accuracy as the aggregate robustness metric.

4.5 Model Setup

We use Gemini-2.5-Flash for primary evaluations. To study cross-architecture generalization, we extend evaluations to Claude Haiku, Sonnet, and Opus; exact API model identifiers are listed in the reproducibility statement. Closed-API behavior may still drift on non-pinned dependencies, which we treat as a reproducibility limitation.

5 Results

5.1 Adversarial Stress Test and Baselines

We compare CDD against baselines to test whether decomposition changes controlled-conflict behavior. The high Claude Standard RAG scores are consistent with Gemini-generated perturbations transferring imperfectly across model families, reinforcing the same-model-bias caveat in §4.1.

Table 1 reports CDD, chain-of-thought prompting (Wei et al., 2022), reflection-token prompting inspired by Self-RAG (Asai et al., 2024), and prompt-only self-filtering inspired by document-filtering RAG (Yoran et al., 2024) on Gemini-2.5-Flash; Figure 3 visualizes the pattern in Appendix B. The reflection baseline is prompt-only, not the trained Self-RAG system.

CDD improves adversarial macro accuracy over Standard RAG (78.1% vs 63.0%). Against Conflict-aware RAG (73.5%), the separable margin concentrates on Entity Swap and Logical Contradiction, where premise isolation is most applicable; Temporal Shift and Distractor Evidence intervals overlap, so we make no robustness claim there.

Average-case calibration. On paired clean/adversarial records, CDD and Standard RAG are statistically tied in harmonic mean (72.23% vs 72.33%, bootstrap $p=0.53$). This hides a trade-off: CDD scores 67.2% clean / 78.1% adversarial, while Standard RAG scores 84.9% / 63.0%. We therefore treat CDD as a conflict-specific diagnostic, and Epi-Scale as a stress test rather than evidence that adversarial contexts remove all answer-relevant information.

Table 1: Adversarial accuracy across perturbation types on the Epi-Scale adversarial records (Gemini-2.5-Flash, N=2,250; ~562–563 samples per perturbation; normal-approximation 95% binomial CIs). Macro Avg. is the unweighted arithmetic mean of the four perturbation accuracies; the rightmost interval propagates the four per-cell binomial uncertainties.

Model & Setting	Entity Swap	Log. Contradict.	Temp. Shift	Distract. Evid.	Macro Avg.
ClosedBook (Zero-shot)	43.7% (± 4.1)	40.7% (± 4.1)	50.0% (± 4.1)	44.4% (± 4.1)	44.7% (± 2.0)
Standard RAG	58.4% (± 4.1)	56.0% (± 4.1)	68.8% (± 3.8)	68.8% (± 3.8)	63.0% (± 2.0)
Vanilla CoT	62.0% (± 4.0)	61.3% (± 4.0)	63.2% (± 4.0)	68.1% (± 3.9)	63.7% (± 2.0)
Reflection Prompting	69.5% (± 3.8)	65.0% (± 3.9)	66.0% (± 3.9)	67.5% (± 3.9)	67.0% (± 1.9)
Prompt-Filtered RAG	68.0% (± 3.9)	64.5% (± 4.0)	65.5% (± 3.9)	67.0% (± 3.9)	66.3% (± 2.0)
Conflict-aware RAG	79.3% (± 3.3)	75.4% (± 3.6)	70.0% (± 3.8)	69.2% (± 3.8)	73.5% (± 1.8)
CDD (Ours)	88.0% (± 2.7)	83.2% (± 3.1)	71.3% (± 3.7)	69.9% (± 3.8)	78.1% (± 1.7)

Table 2: Cross-model replication on the Epi-Scale adversarial records (Claude-family models, N=2,250 adversarial examples). CDD improves adversarial accuracy on all three Claude-family models. The per-model gains are directional point estimates (+2.2/+4.6/+2.6 pp). CDD marginal 95% confidence intervals are shown in the table; paired significance tests for the method differences were not run. This reflects transfer of the accuracy benefit only; trace-perturbation sensitivity transfer is reported as point estimates without paired significance tests (§5.5.3).

Method	Haiku	Sonnet	Opus
Standard RAG	79.0%	76.0%	79.4%
Vanilla CoT	73.6%	73.2%	75.4%
CDD (Ours)	81.2% (± 1.6)	80.6% (± 1.6)	82.0% (± 1.6)

5.2 Ablation Study

Table 3 details component ablations. Removing Premise Isolation (Step 4) drops macro accuracy to 65.1%; removing Divergence Check (Step 3) gives 66.0%. The omitted Temporal Shift/Distractor Evidence mean implied by the table is below Standard RAG, so decomposition can be harmful unless divergence checking and premise isolation are present.

The length-matched sham prompt reaches 40.1% macro (Table 3), below full CDD. Thus the decomposition type matters; generic multi-step structure is insufficient, though parse-failure analysis is needed before attributing the full sham drop to reasoning rather than format fragility.

5.3 Diagnostics & Compute Tradeoff (CDD- α)

For CDD- α , we set $\tau = 0.7$ by inspecting a 100-sample contradiction-score histogram from the adversarial evaluation pool; this is not a separate development set or full sweep. The gate routes 30.0% of samples to CDD and reaches 68.5% adversarial accuracy (Figure 2b), about 1.0 pp above the ana-

Table 3: Component-wise ablation on Gemini-2.5-Flash. Entity Sw. and Log. Con. are shown because they are the localized factual-conflict families where CDD’s strongest baseline margins occur in Table 1. Overall is the unweighted macro average over all four perturbation types (Entity Swap, Logical Contradiction, Temporal Shift, Distractor Evidence). The omitted Temporal Shift and Distractor Evidence cells are not individually recoverable from the archived aggregate table; their two-cell mean is implied by the Overall column.

Ablation Variant	Entity Sw.	Log. Contradict.	Overall
Full CDD	88.0%	83.2%	78.1%
Length-Matched Sham CoT	42.4%	32.0%	40.1%
w/o Step 4 (Isolation)	75.2%	72.0%	65.1%
w/o Step 3 (Diverge)	77.0%	73.5%	66.0%

lytical 30%-random-routing expectation of 67.5%. We present this as an illustrative operating point, not a demonstrated compute/accuracy win.

5.4 Finding 1: Context Compliance Is Measurable

Setup: We sampled 500 TruthfulQA instances and supplied the common human misconception as retrieved context. This is an upper-bound compliance test, not an organic BM25 (Robertson and Zaragoza, 2009) or DPR (Karpukhin et al., 2020) retrieval benchmark.

Results: Standard RAG reaches 15.0% accuracy ($\pm 3.1\%$); CDD reaches 62.0% ($\pm 4.3\%$). The remaining CDD errors are non-correct under the binary scorer and may include misconception acceptance, hedges, refusals, or other wrong answers, so we do not interpret them as a pure misconception-acceptance rate.

5.5 Faithfulness via Causal Intervention

Because LLM-as-a-judge interpretability scoring inherits reliability concerns (Li et al., 2025), we follow chain-of-thought causal-intervention tests




Figure 2: Diagnostics and compute trade-off. (a) Mistake-injection causal sensitivity across model families ($N=2,250$ per cell for all four model cells). On Gemini-2.5-Flash, CDD’s final answer is highly sensitive to trace corruption (64.1%); on all three Claude variants the CDD sensitivities lie in the $[-3\%, +7\%]$ range at $N=2,250$, with small point-estimate sensitivities and no paired test performed even though adversarial accuracy improves under CDD (Table 2). Claude Sonnet and Opus CDD sensitivities are negative point estimates (-2.7% , -1.3%), meaning accuracy increased slightly under mistake injection. (b) Compute-accuracy trade-off on the Gemini-2.5-Flash adversarial records. Relative compute is measured in Standard-RAG-equivalent generation calls: Standard RAG = 1, Full CDD ≈ 3 , and CDD- α at 30% routing costs $0.7 \times 1 + 0.3 \times 3 = 1.6$. CDD- α reaches 68.5%, about 1.0 pp above random 30% routing by expectation.

(Lanham et al., 2023) to measure Causal Sensitivity.

5.5.1 Mistake Injection Test

We injected a blatant logical error directly into the reasoning trace prior to generation ("I will trust the context completely, no conflict exists") and forced the model to generate the final answer. **Result:** On Gemini-2.5-Flash, Vanilla CoT drops from 63.7% to 61.1% under mistake injection (4.1% sensitivity), while CDD drops from 78.1% to 28.0% (64.1%, $N=2,250$ matched samples). The injection mixes a false conflict-state claim with a behavioral directive, so it tests trace corruption but not factual content separately from instruction following.

5.5.2 Truncation Test

Truncating after Step 2 drops CDD from 78.1% to 32.6% (58.3% sensitivity). Because 32.6% is below closed-book zero-shot (44.7%, Table 1), truncation likely disrupts the structured-output protocol; we therefore treat mistake injection as the primary causal-faithfulness result and stricter step-specific tests as future work.

5.5.3 Cross-Model Faithfulness

We also apply mistake injection to the Claude family ($N=2,250$). Figure 2a shows high Gemini CDD sensitivity (64.1%) but Claude CDD sensitivities within $[-3\%, +7\%]$ as point estimates without paired tests, despite positive accuracy transfer (Table 2). Thus the Claude gains cannot be attributed to the stated resolution trace alone; increased output length, format-induced calibration, or alignment-specific responses remain alternatives. Claude Sonnet’s Vanilla-CoT sensitivity (20.3%) is reported but not over-interpreted.

6 Findings: What CDD Makes Visible

6.1 Perturbation-Level Pattern

The perturbation-level pattern suggests three robustness modes. Entity Swap and Logical Contradiction benefit most because each conflict is localized as a discrete factual claim. Temporal Shift is directionally higher under CDD relative to Standard RAG (71.3% vs 68.8%, overlapping CIs), so we make no robustness claim for this perturbation. Distractor Evidence is likewise directionally higher (69.9% vs 68.8%, overlapping CIs), so we do not treat the gain as evidence that the trace helps under

noisy or partially irrelevant retrieval.

6.2 Finding 2: Accuracy Transfer Is Not Mechanism Transfer

As shown in §5.5.3, accuracy transfer (Table 2) and trace-perturbation sensitivity (Figure 2) dissociate: the adversarial-accuracy benefit of explicit decomposition transfers across the Claude family, whereas mistake-injection sensitivity shows no measurable coupling at $N=2,250$ (CDD sensitivities within $[-3\%, +7\%]$) outside Gemini-2.5-Flash. We therefore phrase the finding conservatively—what transfers is the accuracy benefit, not a demonstrated causal coupling between the resolution trace and the final answer. This makes CDD a cross-family accuracy intervention in our controlled setting while leaving its mechanism family-specific; the absence of measurable coupling is consistent with, but does not prove, the absence of a causal footprint.

6.3 Finding 3: Explicit Decomposition Improves Controlled-Conflict Robustness

As reported in Table 1, the clearest controlled-conflict robustness gains occur on Entity Swap and Logical Contradiction, the two localized factual-conflict settings where explicit premise isolation is most directly applicable. This is where CDD separates most clearly from the strongest baseline, Conflict-aware RAG; the table provides the exact values and intervals.

6.4 Future Work

Next evaluations are organic BM25/DPR retrieval for TruthfulQA; multi-document retrieval with one false passage among k true passages; factual-only, rationale-swap, and step-specific interventions; and open-weight replication with logit-level prior/posterior divergence. Potential future directions outside the controlled conflict setting include semi-structured graph/text retrieval and adaptive fusion/reranking settings (Tao et al., 2026). On the efficiency side, long-context compression and reasoning-pruning methods are methodologically relevant to reducing the overhead of decomposition-style diagnostics (Gao et al., 2025; Jiang et al., 2026).

7 Conclusion

We used CDD as an inference-time probe for how standard RAG handles epistemic conflict. Final-answer accuracy does not fully capture conflict

handling: Standard RAG can enter a measurable context-compliance regime under worst-case misconception injection, explicit decomposition can improve controlled-conflict robustness, and accuracy gains can transfer across model families even when trace-perturbation sensitivity between the stated resolution trace and final answer does not.

Equally important, the results delimit when CDD should not be read as a general answer-quality improvement. The clean/adversarial calibration tie and the overlapping intervals on Temporal Shift and Distractor Evidence mean CDD is best interpreted as a diagnostic for cases where retrieved evidence conflicts with stable parametric knowledge. In settings with ordinary retrieval noise, partial relevance, or temporally ambiguous updates, our measurements support only a narrower claim: explicit decomposition can reveal and sometimes improve arbitration behavior, but does not replace retrieval quality control or broader faithfulness tests.

These findings position context compliance as a structural RAG behavior requiring systematic diagnostics. A correct answer can hide whether the model rejected a false retrieval, ignored retrieval, or produced a post-hoc rationale; CDD exposes that arbitration by separating contextual belief, parametric belief, premise isolation, and resolution, then testing whether the final answer is sensitive to trace corruption. This trace-perturbation sensitivity is large on Gemini-2.5-Flash; on the Claude family the same trace improves adversarial accuracy without a large point-estimate causal footprint (Gemini 64.1% vs Claude within $[-3\%, +7\%]$), with paired significance not established, so the diagnostic should be read as model-specific rather than universal. We will release the Epi-Scale split (2,250 clean + 2,250 adversarial records), prompt templates, evaluation settings, and scoring specifications upon publication; the internal experimental codebase is not released.

Limitations

- **Cross-Family Generalization:** The clearest causal-sensitivity signal is on Gemini-2.5-Flash (64.1% mistake-injection sensitivity), with effect magnitude varying across model families. Characterizing the Claude-family behavior—separating architecture, alignment-training, and prompt-sensitivity contributions—is a natural open-weight extension.
- **Causal-Coupling Uncertainty:** On the Claude family, mistake-injection sensitivities lie within $[-3\%, +7\%]$ at $N=2,250$, versus 64.1% on Gemini-2.5-Flash. We report these conservatively as point estimates and read them as no measurable coupling in these runs; paired significance, confidence intervals, or an equivalence/power analysis would be needed to move from this reading to a positive claim about the presence or absence of a causal footprint.
- **TruthfulQA Closed-Book Control:** The 15.0% figure characterizes context-injected behavior under our binary scorer; adding a no-context control on TruthfulQA-500 would further quantify the prior-versus-context gap and is a direct extension.
- **Baseline Prompt Asymmetry:** Because the Standard RAG prompt implicitly encourages context-reliance under adversarial context while CDD instructs explicit conflict checking, we include a conflict-aware prompt baseline to separate decomposition from instruction effects. CDD retains a margin on the localized factual conflicts (Entity Swap, Logical Contradiction) and the margin closes on Temporal Shift and Distractor Evidence; we accordingly read part of the overall gain as attributable to the conflict-checking instruction and part to decomposition.
- **CDD- α Threshold Tuning:** We select the CDD- α threshold $\tau = 0.7$ from a 100-sample histogram inspection on the adversarial evaluation pool rather than a separate development set or full sweep; a $\tau \in \{0.3, 0.5, 0.7, 0.9\}$ ablation with a separately fixed gate would be needed before treating the Pareto point as deployment-relevant. We present CDD- α as an illustrative compute/accuracy operating point rather than a tuned deployment setting.
- **TruthfulQA Realism:** Using the explicit top misconception is a deliberate worst-case probe; organic BM25 retrieval would measure average-case degradation and is the natural follow-up.
- **Prior Answer Not Context-Isolated:** In the CDD protocol, Step 2 elicits the parametric answer after the context is visible, so it is a prior reported under the CDD prompt rather than a context-free parametric answer. We therefore treat the belief-revision framing as conceptual; a separate no-context elicitation would tighten the prior/posterior separation and is a direct extension.
- **Conceptual-Only Belief-Revision Framing:** Our belief-revision framing is conceptual rather than a token-level divergence (e.g., JSD); a white-box open-weight instantiation would enable direct quantitative validation of the compliance/resolution distinction.
- **Stricter Faithfulness Tests:** Because mistake injection combines factual content with a behavioral directive, a sharper causal picture would come from factual-only Step-2 corruption, rationale swaps, step-specific corruption, and answer-hidden interventions with contexts that do not reveal the answer; we outline these as the immediate next interventions.
- **Closed-API Drift:** Closed-API behavior may drift under provider updates, and evaluation configuration choices can change benchmark verdicts (Li et al., 2026b); version-pinned open-weight replication would remove one major source of variance.
- **Single-Run Closed-API Evaluation:** The reported closed-API results are single-run evaluations at temperature 0.0. Temperature 0.0 does not guarantee deterministic behavior for hosted APIs, so small directional gaps, especially in the Claude-family replication, should be interpreted with caution until repeated runs or open-weight replications are available.

Reproducibility Statement

Due to API-provider terms and maintenance constraints, we do not release the full internal experimental codebase, which contains non-generalized

infrastructure and API-specific wrappers. We will release the Epi-Scale split (2,250 clean + 2,250 adversarial records), prompt templates, evaluation settings, and scoring specifications upon publication; the internal experimental codebase is not released. The release split contains 2,250 clean-context records and 2,250 adversarial-context records, balanced across HotpotQA, Natural Questions, and FEVER and across the four perturbation types. The API identifiers used in the reported experiments are:

Gemini	<code>gemini-2.5-flash</code>
Claude Haiku	<code>claude-haiku-4-5-20251001</code>
Claude Sonnet	<code>claude-sonnet-4-6</code>
Claude Opus	<code>claude-opus-4-6</code>

The Gemini endpoint has an announced discontinuation date of 2026-06-17, so exact replication against the same endpoint may no longer be possible. The Haiku ID is a dated snapshot. From the Claude 4.6 generation onward, a dateless ID such as `claude-sonnet-4-6` is a version-pinned snapshot rather than an evergreen alias, so the model weights for a given ID are fixed; however, the surrounding serving infrastructure (e.g., request routing, moderation classifiers, and sampling) may change over time, which we treat as a reproducibility limitation. Experiments were conducted between January and March 2026. Closed-API model behavior may drift on non-pinned dependencies (e.g., moderation layers), which we treat as a reproducibility limitation. All runs use temperature 0.0 and greedy decoding constraints. CDD- α uses a prompted Gemini-2.5-Flash contradiction gate rather than a separate NLI checkpoint; the gate returns a parsed scalar score from a prompt asking whether the context contradicts the hypothesized answer. Dataset construction draws evenly from HotpotQA, Natural Questions, and FEVER; perturbations are generated along four mutually exclusive axes and audited manually on 50 examples per subgroup (92% valid conflict generation rate). Metrics use normalized string matching with lowercase, punctuation stripping, and alias mapping; FEVER labels are mapped to Boolean support/refute classes in the appendix. Confidence interval computation follows the formulas detailed in §4.3. Compute resources: All experiments use closed-API inference at standard rate limits; we did not perform local fine-tuning or training.

References

- Carlos E. Alchourrón, Peter Gärdenfors, and David Makinson. 1985. On the logic of theory change: Partial meet contraction and revision functions. *The Journal of Symbolic Logic*, 50(2):510–530.
- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. 2024. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations*.
- Hung-Ting Chen, Michael J.Q. Zhang, and Eunsol Choi. 2022. Rich knowledge sources bring complex knowledge conflicts: Recalibrating models to reflect conflicting evidence. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 2292–2307, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.
- Zhiyuan Cheng and Longying Lai. 2026. [Energy-efficient on-device RAG on a mobile NPU: System design and benchmark on snapdragon x elite](#). *Preprint*, arXiv:2606.11257.
- Zhiyuan Cheng, Longying Lai, and Yue Liu. 2026a. [Sustainable hybrid document-routed retrieval for financial RAG: Resolving the robustness-precision trade-off](#). *Preprint*, arXiv:2603.26815.
- Zhiyuan Cheng, Longying Lai, Yue Liu, Kai Cheng, and Xiaoxi Qi. 2026b. [Enhancing financial report question-answering: A retrieval-augmented generation system with reranking analysis](#). *Preprint*, arXiv:2603.16877. Accepted at ICECET 2026.
- Zhichao Fan, Zexin Zhuang, and Yanhang Li. 2026. [Probe choice changes canary-memorization verdicts: Three post-hoc disagreement case studies in a text-dominant LoRA-tuned autoregressive testbed](#). In *ICML 2026 Workshop on Foundations of Deep Generative Models: Understanding Memorization, Generalization, and Reasoning*.
- Yaxin Gao, Yao Lu, Zongfei Zhang, Jiaqi Nie, Shanqing Yu, and Qi Xuan. 2025. [DSPC: Dual-stage progressive compression framework for efficient long-context reasoning](#). *Preprint*, arXiv:2509.13723, arXiv:2509.13723.
- Yuxuan Jiang, Dawei Li, and Francis Ferraro. 2026. [DRP: Distilled reasoning pruning with mathematical skill-aware step decomposition for efficient large reasoning models](#). In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 4020–4039, San Diego, California, United States. Association for Computational Linguistics.
- Vladimir Karpukhin, Barlas Oguz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 6769–6781, Online. Association for Computational Linguistics.

- Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. 2019. Natural questions: A benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466.
- Tamera Lanham, Anna Chen, Ansh Radhakrishnan, Benoit Steiner, Carson Denison, Danny Hernandez, Dustin Li, Esin Durmus, Evan Hubinger, Jackson Kernion, Kamilė Lukošiuūtė, Karina Nguyen, Newton Cheng, Nicholas Joseph, Nicholas Schiefer, Oliver Rausch, Robin Larson, Sam McCandlish, Sandipan Kundu, and 11 others. 2023. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems*.
- Dawei Li, Bohan Jiang, Liangjie Huang, Alimohammad Beigi, Chengshuai Zhao, Zhen Tan, Amrita Bhat-tacharjee, Yuxuan Jiang, Canyu Chen, Tianhao Wu, Kai Shu, Lu Cheng, and Huan Liu. 2025. [From generation to judgment: Opportunities and challenges of LLM-as-a-judge](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 2757–2791, Suzhou, China. Association for Computational Linguistics.
- Yanhang Li, Zhichao Fan, and Zexin Zhuang. 2026a. [Auditing reasoning-trace memorization claims after unlearning with head-conditioned canaries](#). *Preprint*, arXiv:2605.18891.
- Yanhang Li, Zhichao Fan, and Zexin Zhuang. 2026b. [SafetyRepro: Configuration-conditional rank instability on alignment benchmarks](#). *Preprint*, arXiv:2605.25492.
- Stephanie Lin, Jacob Hilton, and Owain Evans. 2022. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3214–3252, Dublin, Ireland. Association for Computational Linguistics.
- Shayne Longpre, Kartik Perisetla, Anthony Chen, Nikhil Ramesh, Chris DuBois, and Sameer Singh. 2021. Entity-based knowledge conflicts in question answering. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7052–7063, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das, Daniel Khashabi, and Hannaneh Hajishirzi. 2023. When not to trust language models: Investigating effectiveness of parametric and non-parametric memories. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 9802–9822, Toronto, Canada. Association for Computational Linguistics.
- Hrishikesh Paranjape, Abhishek Mandal, and Xian Sun. 2026. [Optimality of sequential filtering under independent cost and selectivity models](#). In *2026 IEEE International Conference on Electro/Information Technology (EIT)*. IEEE. Accepted at EIT 2026.
- Stephen Robertson and Hugo Zaragoza. 2009. The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval*, 3(4):333–389.
- Weijia Shi, Xiaochuang Han, Mike Lewis, Yulia Tsvetkov, Luke Zettlemoyer, and Scott Wen-tau Yih. 2024. Trusting your evidence: Hallucinate less with context-aware decoding. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pages 783–791, Mexico City, Mexico. Association for Computational Linguistics.
- Xian Sun, Wei Gao, Yingshuo Wang, Lingdong Kong, Yanhang Li, Zhichao Fan, Zexin Zhuang, Wenlong Dong, Zhiyuan Zheng, Hrishikesh Paranjape, Abhishek Mandal, and Johnny R. Zhang. 2026. [Beyond accuracy: Measuring bias acknowledgment in chain-of-thought reasoning for responsible AI evaluation](#). *Preprint*, arXiv:2606.15127. ICML 2026 Workshop on Trustworthy AI for Good.
- Yicheng Tao, Yiqun Wang, Xiangchen Song, Xin Luo, Kai Liu, and Jie Liu. 2026. [GRASP: Plan-guided graph retrieval with adaptive fusion and reranking on semi-structured knowledge bases](#). *arXiv preprint arXiv:2605.30237*.
- James Thorne, Andreas Vlachos, Christos Christodoulopoulos, and Arpit Mittal. 2018. FEVER: A large-scale dataset for fact extraction and VERification. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 809–819, New Orleans, Louisiana. Association for Computational Linguistics.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel Bowman. 2023. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. In *Advances in Neural Information Processing Systems*.
- Fei Wang, Xingchen Wan, Ruoxi Sun, Jiefeng Chen, and Sercan Ö. Arık. 2025. Astute RAG: Overcoming imperfect retrieval augmentation and knowledge conflicts for large language models. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*,

pages 30553–30571, Vienna, Austria. Association for Computational Linguistics.

Yingshuo Wang, Xian Sun, Lingdong Kong, Wei Gao, Yanhang Li, Zhichao Fan, and Zexin Zhuang. 2026. [Do time series foundation model benchmarks hide regime-dependent failures? evidence from traffic speed forecasting](#). *Preprint*, arXiv:2606.18367. Accepted at the Workshop on Forecasting as a New Frontier of Intelligence, ICML 2026.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. 2022. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*.

Kevin Wu, Eric Wu, and James Zou. 2024. ClashEval: Quantifying the tug-of-war between an LLM’s internal prior and external evidence. *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*.

Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, and Yu Su. 2024. Adaptive chameleon or stubborn sloth: Revealing the behavior of large language models in knowledge conflicts. In *International Conference on Learning Representations*.

Rongwu Xu, Zehan Qi, Zhijiang Guo, Cunxiang Wang, Hongru Wang, Yue Zhang, and Wei Xu. 2024. Knowledge conflicts for LLMs: A survey. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 8541–8565, Miami, Florida, USA. Association for Computational Linguistics.

Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. 2024. Corrective retrieval augmented generation. *arXiv preprint arXiv:2401.15884*.

Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. 2018. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2369–2380, Brussels, Belgium. Association for Computational Linguistics.

Ori Yoran, Tomer Wolfson, Ori Ram, and Jonathan Berant. 2024. Making retrieval-augmented language models robust to irrelevant context. In *International Conference on Learning Representations*.

Zexin Zhuang, Yanhang Li, and Zhichao Fan. 2026. [Pre-registering the detectable effect: A paired-MDE budget for 4-bit quantization benchmarks, with a pilot audit](#). *Preprint*, arXiv:2605.28873.

Wei Zou, Runpeng Geng, Binghui Wang, and Jinyuan Jia. 2025. PoisonedRAG: Knowledge corruption attacks to retrieval-augmented generation of large language models. In *34th USENIX Security Symposium (USENIX Security 25)*, pages 3827–3844. USENIX Association.

A Appendix: Reproducibility Details

All evaluations use temperature 0.0 and greedy decoding constraints. String normalization mapping for FEVER: [supports, true, yes] → True; [refutes, false, no, contradicts] → False. **Standard**

RAG Template:

```
Context: {context}
Question: {question}
Based on the context, answer
concisely.
```

ClosedBook Template:

```
Question: {question}
Answer concisely.
```

Vanilla CoT Template:

```
Context: {context}
Question: {question}
Think step-by-step, then answer
the question. Wrap final answer
in <final_answer>.
```

Reflection Prompting Template:

```
Given the context, answer the
question. You must include
reflection tokens.
Context: {context}
Question: {question}
Format your response as:
[Relevant: Yes/No]
[Supported: Fully/Partially/No]
[Contradicts_Prior: Yes/No]
Final Answer: <your concise
answer>
```

Prompt-Filtered RAG Template:

```
Context: {context}
Paraphrase this context.
Critically, REMOVE any factual
claims that severely violate
established world knowledge or
history. If a claim seems fake,
omit it.
Filtered Context:
{filtered_context}
Question: {question}
Answer concisely based on the
filtered context and your world
knowledge. Final Answer:
```

Conflict-Aware RAG Template:

```
Context: {context}
Question: {question}
Carefully answer the question
using both the retrieved context
and your own knowledge.
If the retrieved context is
consistent with well-established
facts, use it as the primary
evidence.
If the retrieved context appears
to conflict with
well-established scientific or
historical knowledge, explicitly
prioritize the more reliable
information instead of blindly
following the retrieved context.
Do not generate intermediate
reasoning.
Answer concisely.
Wrap your final answer in
<final_answer> tags.
```

Length-Matched Sham CoT Template:

```
Context: {context}
Question: {question}
Follow these five steps before
answering.
Step 1: Restate the question in
your own words.
Step 2: Restate the retrieved
context in your own words.
Step 3: List the named entities
or key terms mentioned in the
question and context.
Step 4: Identify the broad
question type (e.g., person,
place, date, fact verification,
or explanation).
Step 5: Answer the question
concisely based on the context.
Wrap your final concise answer
in <final_answer> tags.
```

CDD Template:

```
You are resolving cognitive
conflicts in retrieval-augmented
generation.
Context: {context}
Question: {question}
Follow these steps:
Step 1: Extract the answer
implied by the context.
Step 2: State your
parametric/internal knowledge
answer to the question.
Step 3: Compare the two answers.
Do they conflict?
Step 4: If they conflict,
isolate the premises causing the
conflict. Evaluate if they
violate established
history/science.
Step 5: Resolve the conflict
and provide the most factually
reliable final answer.
```




Figure 3: Adversarial accuracy on the Epi-Scale adversarial records (Gemini-2.5-Flash, N=2,250; ~562–563 examples per perturbation; normal-approximation 95% binomial CIs). CDD scores 88.0%, 83.2%, 71.3%, and 69.9% across the four perturbation types. CDD is highest numerically in each column, with the statistically clearest separation on Entity Swap and Logical Contradiction; the Temporal Shift and Distractor Evidence gains are directional only (overlapping CIs).

Wrap your final concise answer
in <final_answer> tags.

CDD- α Prompted Contradiction Gate:

```
Premise: {context}
Hypothesis: The answer to
'{question}' is
{parametric_answer}.
Does the Premise explicitly
contradict the Hypothesis? Give
a probability from 0.0 to 1.0.
```

spot check over 200 records total; inter-annotator agreement was not measured, and per-record validity labels were not retained for a filtered-subset sensitivity check.

B Appendix: Epi-Scale Details

Epi-Scale comprises 2,250 base examples drawn evenly from HotpotQA (Yang et al., 2018), Natural Questions (Kwiatkowski et al., 2019), and FEVER (Thorne et al., 2018). Each base example is represented as one clean-context record and one adversarial-context record, yielding 4,500 evaluation records. The adversarial records are divided uniformly among Entity Swap, Temporal Shift, Logical Contradiction, and Distractor Evidence. The perturbations were generated using gemini-2.5-flash with prompt templates constraining semantic alterations to exactly one variable axis. To ensure the LLM-based perturbation engine successfully generated genuine conflicts without destroying syntax, we manually audited a random sample of 50 adversarial records per subgroup, finding a 92% valid conflict generation rate. The audit was conducted as a manual